

The agentic layer

How model-driven assistance attaches to the Finné dispute system without touching the money path, the decision, or the record. Extends PRD v1.1 (Finné on Circle Refund Protocol). The v1 spine is unchanged; every agent in this addendum attaches from the side and can be removed without breaking the loop.

Version 1.0 · 3 August 2026 · Owner: Arko Ganguli (CEO) · Engineering: Abhishek Sira Chandrashekar (CTO) · Status: approved for build · Hackathon slice ships with the 9 August submission; the rest is post-hackathon roadmap

A. Purpose and relationship to PRD v1.1

PRD v1.1 ships a deterministic dispute system: chain events decoded by code, seven pure-function checks, a findings-only brief, a human decision with a written reason, and money that moves only when the reviewer signs. This addendum specifies where models are added to that system, what each one is permitted to produce, and how the system behaves when any of them fails or is switched off.

The design rule throughout: **agents attach from the side, never on the spine**. Each agent pre-fills, stamps, orders, drafts or proposes. None classifies without confirmation, none writes a fact without a source hash, none renders a verdict, none holds keys, and none is required for the loop to complete. Removing every agent in this document restores v1 exactly.

B. Principles, extended

P1 to P5 of the PRD stand unmodified. This addendum adds five.

P6 · Attention, not conclusion. Model output for the reviewer directs attention: the questions the case turns on, what each outcome requires, what remains unresolved. It never marks an outcome as correct, in words or by framing. Frames are generated with the same structure and depth for both sides, and this symmetry is tested, not assumed.

P7 · Stamped or silent. Every model claim carries a pointer to a hashed source already in the evidence record. A claim that cannot be stamped is not shown. The reviewer trusts the stamp, not the model.

P8 · Degrade to v1. Every model call is optional. Each station has a defined degraded path that is the v1 behaviour, and CI proves the full demo loop passes with all models unplugged. A hard timeout applies to every call; on failure the station degrades silently.

P9 · Data stays home. Inference runs on Finné-controlled machines with open weights. No case content, evidence, terms or party detail is sent to an external model API. For the hackathon this is a local Ollama runtime at zero cost; the same posture becomes the no-external-model tier in the merchant offer.

P10 · Libraries grow; models never decide. New decoders, reader types, policy clauses and law lines are authored offline with model assistance, reviewed by a person, versioned, and only then used at runtime. No agent gains a capability mid-case.

C. Agent catalogue

Six agents, each named against the PRD component it attaches to. Phase marks what ships for 9 August.

ID	AGENT	ATTACHES TO	PRODUCES	DEGRADED PATH	PHASE
----	-------	-------------	----------	---------------	-------

A1	Intake classifier	C4 dispute endpoint	Proposed claim type, contested amount and evidence checklist, pre-filled into the dispute form. The human confirms or overrides; the stored claim type is always human-confirmed.	Blank form; party selects the claim type (v1)	POST-HACKATHON
A2	Reader fan-out	C5, triggers T2/T3	A planner dispatches one reader per messy source (PDF, thread, link) from a fixed catalogue, one level deep. Each returns a summary stamped with the sha-256 of the source it read, written as annotations on existing EvidenceItems, never as new facts.	Evidence renders raw; reviewer reads it unaided (v1)	POST-HACKATHON
A3	Relevance selector	C5 brief assembly	An ordering and highlight over evidence and findings, keyed to the claim type. Presentation only: the full set stays in the case room and the reviewer can expand past the selection. Nothing is filtered out.	Chronological order (v1)	POST-HACKATHON
A4	Decision frame	Screen 4, beside the reason box	Turning questions, outcome requirements and unresolved items, per section E. Verdict-free by construction. Reviewer may accept a line into the reason box, edit it, or discard the frame.	Empty reason box; reviewer writes unaided (v1)	HACKATHON
A5	Anomaly watcher	C3/C5 boundary, new trigger T0	Proposed cases from patterns in decoded events and receipts: payments with no matching work order, repeat clawbacks against one recipient, clustered reserve drawdowns. Proposal only; a person at the platform opens the case.	Cases open on complaint only (v1)	POST-HACKATHON
A6	Offline authors	Outside the runtime	Model-assisted drafts of new decoders, reader types, check functions, policy clauses and law lines, merged to the catalogue only after human review. This is the sole route by which agents gain capability.	Hand-authored entries	CONTINUOUS

Invariants across all six. No agent process holds key material; the FIN-45 startup assertion and the lint rule barring import paths to signing code extend to every agent service. Every model call logs model name, version, input hash, output and the human's eventual action. The money path is untouched: USDC moves only when the platform wallet pays or the reviewer's wallet refunds.

D. Hackathon slice

Four things ship by 9 August, all inside the existing milestone plan and none on the critical path.

1. **Local model runtime.** Ollama serving gpt-oss-20b as primary (Apache 2.0, runs on the build laptop), Mistral Small 4 as the vision-capable alternative. Zero cost, no API key, no external call. This replaces stretch S3's external call and upgrades its demo line: the model runs on our machine and no case data leaves it.
2. **The policy pack.** The platform's payout terms seeded as a hashed EvidenceItem, three clauses plus a one-line governing-law note (section F). The rulebook the case is judged under becomes tamper-evident evidence.
3. **Two clause checks.** Grace window and acceptance status join the deterministic engine as checks eight and nine (section F). Pure functions over timestamps and clause parameters.
4. **The decision frame.** A4 as specified in section E, rendered on Screen 4, filled largely by templates and computation, with the local model phrasing only the turning questions.

A1, A2, A3 and A5 are post-hackathon. Their schemas are defined now (section H) so the corpus logging starts on day one, but no runtime code ships for them in this build.

E. The decision frame

The frame is the replacement for a draft verdict. It structures the reviewer's decision without taking a side in it, and most of it is not a model.

E.1 Content, three parts

Turning questions. Derived from failed or contested checks and the clauses they invoke. Example: was video three delivered on time under clause 4? The delivery message is timestamped 2 June; the due date was 1 June; clause 4 allows 48 hours. *Source: the model phrases these from the findings table. It poses questions; it never answers them.*

Outcome requirements. For each of the four actions on Screen 4, what would have to be true. Example: refund requires a written rejection within the clause 4 window, or grounds under the clause 9 fraud exception. *Source: authored templates per claim type, filled from findings by code. No model involvement, which is what makes the outcome lines safe to name outcomes.*

Unresolved items. Gaps computed from the record: an unanswered reply, uncountered evidence, a contested amount that does not match a tranche, absent acceptance criteria. *Source: database checks. No model.*

E.2 Symmetry rule

The frame is generated blind to which outcome the findings favour: same structure, same depth, for both parties' positions, every time. The test is mechanical: every outcome path enumerated, every party's material cited at equal depth, and a snapshot test comparing frame structure across mirrored seeded cases. A frame that enriches one side is a defect of the covert-steer class and blocks release.

E.3 Placement and record

The frame renders on Screen 4 only, beside the mandatory reason box, and is stored as a `DraftFrame` record separate from the `AgentBrief`. The brief schema keeps no verdict field and no frame content; P1's letter is preserved: briefs carry findings, frames live with the decider. Reviewer actions: accept a line into the reason box, edit, or discard. The action taken is logged per line for the corpus.

E.4 What the frame is not

Not a recommendation, open or disguised. Ordering evidence toward an outcome, phrasing one side stronger, or burying passed checks are covert steering and are treated as defects, not features. An open recommendation, if ever shipped, is a separately decided later stage with its own decision-log entry and design-partner sign-off; it is out of scope here.

F. The standard layer: terms and law

Facts alone decide nothing; the reviewer needs what was owed and what is allowed. The standard enters as evidence, not as model output.

F.1 Policy pack

One seeded document, the platform's payout terms, uploaded and hashed like any EvidenceItem, cited by clause number throughout the case. Demo clauses:

- **Clause 4.** A deliverable submitted within 48 hours after its due date is treated as on time unless the platform rejects it in writing within that window.
- **Clause 7.** A deliverable not disputed within 14 days of submission is deemed accepted; accepted work is payable.
- **Clause 9.** The platform may seek refund of a payout only before the deliverable is deemed accepted, except for fraud.

F.2 Law line

The pack carries a governing-law note: the contract is governed by Irish law, and deemed-acceptance clauses of this kind are enforceable between businesses; a clawback after acceptance is a contractual claim requiring grounds, not a self-help right. One jurisdiction, one authored sentence, curated offline by the legal founder, retrieved in the case, never generated at runtime. This is A6 in miniature and the seed of the law library.

F.3 Checks eight and nine

CHECK	LOGIC	OUTPUT
Grace window	Submission timestamp minus due date, compared against the clause 4 window; written rejection searched for within the window.	pass fail per deliverable, with dates and the clause cited
Acceptance status	Days between submission and dispute filing, compared against the clause 7 period.	deemed-accepted date, dispute date, pass fail with clause cited

Both are pure functions over timestamps and clause parameters, same shape as the existing seven, tested the same way. The demo scenario resolves as: one day late but inside clause 4; deemed accepted 16 June under clause 7; dispute filed 20 June; no clause 9 ground raised.

G. Model runtime

ITEM	SPECIFICATION
Serving	Ollama on the build machine; model pulled once, pinned by digest, recorded in docs/models.md.
Primary model	gpt-oss-20b, Apache 2.0, text-only, sufficient for question phrasing and the narrative summary.

Alternative	Mistral Small 4, Apache 2.0, vision-capable; adopted when A2's document readers arrive, usable now if preferred.
Budget	5-second hard timeout per call, inherited from S3. On timeout or malformed output the station degrades per P8.
Cost	Zero. No API key exists in any environment; the secret-scan and FIN-11 assertions extend to model keys.
Output filter	The S3 post-filter (no outcome words in model-generated text) applies to turning questions and the narrative. Outcome-requirement lines are template-authored and exempt by construction.
Swap rule	The contract is the schema, not the model. Any model change is a config change plus a regression run; product code does not reference model names.

H. Schemas and corpus

Every agent writes into a typed record validated before render. Defined now, including for post-hackathon agents, so logging is uniform from day one.

RECORD	REQUIRED FIELDS	BANNED FIELDS
DraftFrame	case ref, questions[] (text, findings refs), requirements[] (outcome, template id, filled params), unresolved[] (kind, refs), model id+digest, generated at	Any verdict, score, confidence or ranking of outcomes
EvidenceAnnotation	evidence ref, source sha-256, summary, spans cited, reader type, model id+digest	New facts without a source hash; references outside the evidence table
ProposedCase	pattern id, events cited, receipt refs, proposal text	Auto-open flag; any case status field
PolicyClause	pack ref, clause number, text, parameters, jurisdiction, author, review ref, version	Runtime-generated entries

Corpus logging. Every model call stores input hash, output, validation result and the human's action (accepted, edited, discarded, overridden). Resolved cases thereby accumulate labelled examples per task, the training set for fine-tuned small models later and the audit trail now.

I. Post-hackathon build order

Ordered by risk each agent carries, lowest first: **A2** readers (annotate hashed evidence only), **A3** selector (orders a list), **A1** classifier (pre-fills a confirmed field), **A5** watcher (proposes cases, tuned against live dispute volume), then any A4 stage two as a separately decided step. A6 runs throughout. Each lands as a sellable upgrade on an already-trusted base; the v1 sales line, no model in the path, remains true for every merchant on the no-model tier.

J. Risks

RISK	ODDS	MITIGATION
Covert steer in the frame	Medium	Outcome lines template-authored; symmetry snapshot test on mirrored cases; per-line action logging surfaces skew; defect class blocks release.

Local model latency on the demo laptop	Medium	5s timeout with silent degrade; frame renders from templates and computation alone if phrasing fails; rehearsal includes a models-unplugged run.
Malformed model output	High, tolerable	Schema validation before render; failures degrade, never display.
Scope creep into post-hackathon agents	High	Only sections D items ship; schemas exist, runtimes do not; cut order in the issue list.
Clause pack read as legal advice	Low	Fictional demo terms; law line authored and attributed; disclaimer beside the pack in the case room.

K. Decisions log, additions

ID	DECISION	STATUS
D7	All model inference runs on Finn�-controlled machines with open weights; no external model API in any environment. Hackathon runtime is Ollama with gpt-oss-20b primary, Mistral Small 4 alternative.	Decided 3 Aug
D8	A4 ships as a decision frame (attention, not conclusion), not a recommendation. Any open recommendation is a separately decided later stage.	Decided 3 Aug
D9	Hackathon slice is section D exactly: local runtime, policy pack, checks eight and nine, decision frame. A1, A2, A3, A5 are schema-only in this build.	Decided 3 Aug
D10	Demo policy pack governed by Irish law with one authored law line; law entries are curated offline, never generated at runtime.	Decided 3 Aug
D11	Corpus logging (input hash, output, validation, human action) starts with the first model call and covers every agent task.	Decided 3 Aug